

The Project, in short

- You will work in teams
- to build a distributed system
- doing something **awesome**
- and expanding your knowledge base

The Project

- At least 4 nodes with DS technologies:
 - Communications
 - e.g. protocol transparency, location transparency, source independence, and destination independence
 - Coordination
 - Sharing of some state information to make distributed decisions



The Project: DS Features

- Scope your project to include some of these features:
- Failure handling (node, process, network)
- Multimedia data (audio / video)
 - learn about realtime features
- Bootstrap configuration, but not from a bootstrap server
 - Okay to use DNS solutions

The Project: DS Features (2)

- QoS and Resource management
 - reservations to provide guarantees of performance
- Security
 - online commerce techniques
- Middleware
 - CORBA, Enterprise Java Beans: The power of abstraction in taming DS complexity




The Project: MUST

- Project MUST include some sort of **skills challenge**
 - Something that pushes your group to learn and employ better technology
 - **New programming language (Erlang, Go, Reia, ?)**
 - In general: No Java, C, C++
 - New platform / middleware (CORBA, ?)
 - Research techniques (Content-centric Networking, ?)
 - New hardware platform (Embedded sensor nodes, ?)
 - Android not considered “New”
 - Management/SW process techniques (Agile, ?)

The Project: MUST NOT

- Your project MUST NOT be designed to centralize decision making
 - Don't use a distributed filesystem / database to let all nodes know all information
 - Don't send all updates to a central server where decisions are made
- Also, must not use *a. priori*. knowledge of a bootstrap server's IP address



Scaling Analysis

- Must analyze how big N can be
 - Measure your bandwidth usage and patterns
 - Understand your critical path / section
 - Properly use queuing theory, appropriate analytic (prob/stat?) techniques
- Can you state how much better you are than a centralized solution?
- Simulation can be your friend
 - For P2P projects: PlanetSim, P2PSim, PeerSim, ...

How to get an 'A' on your Project

- Your team built something that:
 - Works great
 - Perhaps has a bug or two, but you have good workarounds in place
 - and you've tested it well enough to be confident in it
 - Fulfills the goals described in your specification
 - You didn't cut out major parts because you couldn't get them working
 - Is challenging
 - Challenged the technology skills of all team members
 - Applies DS principles well
 - Implements fundamental algorithms

How to get an 'A' on your Project

- You:
 - Were a good team member
 - Contributed often
 - Did lots of work
 - Focused on the mission
 - No team member ever complained about you

How to get a 'B' on your Project

- Your team built something that:
 - Mostly works
 - Fulfills some of the goals described in your specification
 - But, you threw out the hard stuff when it didn't work
 - Is somewhat challenging
 - Written in java, using std libraries or stuff grabbed from the net
 - Connects your 4 laptops together, nothing else
 - And you aren't sure how it would scale to other machines
 - Applies networking principles well
 - to send all information to a central server / database for action



How to get a 'B' on your Project

- You:
 - Were an undistinguished team member
 - Contributed
 - Did lots of work
 - But didn't quite get it to execute well
 - But not when your teammates needed you

Some Basic Project Rules

- You can continue to refine specifications during the course of the semester
 - Typical: requirements change all the time!
 - A good design can deal with changes (or enhancements to supported requirements)
 - Use layering
 - lawyers are not so good; but layers are!
 - Talk to TA before making major changes

Some Basic Project Rules

- You can propose your own project
 - get credit for doing that project you have been itching to do!
 - talk to faculty members you know: will likely give you a close guide to work with
- Don't panic ...
 - ... but do start early!
 - We will work with you to get a reasonable project description
 - Reasonable reuse of software from other sources is allowed (foundation classes, algorithms, commercial back-ends, etc.)
 - Re-use of platforms / problems from your research is encouraged

Teams

- Each project team consists of 4 students
- Teams will be assigned
 - Method to this (madness): force you to learn how to talk to other tech folks (yes, just like the real world)
- Manage your group dynamics
 - It's your problem....
 - ...unless you tell me early enough
- Each team will work with a particular TA
 - Meet at least bi-weekly



Team Skills

- Planning
- Expectation-setting
 - Team contracts, ground rules
 - Establish expectations and commitments
- Delegating
 - Know the skills of your teammates

Team Skills (2)

- Leadership
 - Choose leadership style and roles
 - Can be tough in a group of peers
- Teamwork
 - There are a range of roles you will play, all of which contribute towards your team's success

Team Skills

- Conflict Resolution
 - Friction sometimes can be creative
 - Disagreements over technical approach
 - Conflict often comes from communication barriers, so make sure you are listening
 - Can't get personal, or everyone loses

Handling conflict

- Direct approach: confront the issue
- Bargain: find a compromise
 - Especially good for handling conflicting solutions to a problem or design
- Enforcement of team rules: Only as a last resort, as it causes lots of hard feelings
- Retreat: Avoid or work around the problem
- De-emphasis: Move in a new direction, towards areas of agreement

Project Ideas



- **Automotive applications**
 - OpenXC has tools and APIs to develop vehicular applications and access data sources in the car
 - I have a former-student contact at Ford who has offered a development kit to an 842 team
 - openxcplatform.com

Video Management Platforms

- **"Video on Demand" Services**
- **Streaming platforms, esp P2P**
 - Timely delivery of audio/video streams
 - Cached near the point-of-use to minimize/eliminate congestion
- Can you outdo these guys?






- **Exploit Skype Infrastructure**

- Skype is a very successful P2P system
- Has an open API. Leveragable?
- Could you backup your computer to your Skype contacts?

- **P2P Data Analysis**

- To what degree could “big data” benefit from P2P computation?
- Ex: <http://nssdc.gsfc.nasa.gov/> for Astronomy data



- **Distributed Collaboration**

- Audio / video conferencing: Many-to-many
- Collaborative whiteboarding w/audio support
- Distributed chat: support text, graphics, audio, video?
- Minimize BW needs with smart video and audio

- **Distributed "Active Network" Services (SDN)**

- We inject postscript into printers as computer programs
- Similarly, send new routing, transformation, filtering policies to routers and switches within networks
- The whole internet becomes programmable!
- Quality adapts to available BW, display size, etc.



Network and Resource Management

- **Distributed Quality of Service Management**

- QoS tradeoffs have to be made all the time!
- infrastructural support for delivering variable quality of service based on BW availability, fees charged, etc.
- billing and tracking support

- **QoS aware End to End Applications**

- build applications like videoconferencing “automatically” adapt to congestion conditions
- Treat premier customers with higher QoS
- Time-based/bit-based charging services



Infrastructure Support

- **Visualization/Control of Distributed Execution**

- A debugging/monitoring environment for distributed systems
- Log events and resource usage locally
- Log userlevel events and kernellevel events (by emulation if needed)
- Batch and ship events to a remote console to visualize them on a single screen
- Add 2way communications to each machine from monitoring machine to stop/start processes, add/filter events
- Aids performance tuning and visualization



- **Multicast Support**

- Multicast IP and RealTime Protocol (RTP)
- Session Initiation Protocol (SIP)
- synchronize multiple streams, merge multiple streams
- split a merged stream into components


- **A "PushPull" communications architecture**

- may or may not be modeled after the publisher/subscriber
- server can “push” data to subscribed clients on regular basis
- a client can “pull” data from server on demand
- support for different types of data



- **Wireless Sensor Networks**

- Large numbers of physically distributed sensors
 - Including cameras, microphones, bio-sensors, chemical sensors, etc.
- Distributed and centralized/replicated processing of audio/video for “special” events, faces, etc.
- Persistent logging of all events with easy retrieval
- Flexible querying of monitored data (“where is Bill now?”)
- Levels of access control
- Flexible front-ends
- Remote administration



- **Smart Electronics for the Home**
 - Wireless/wired extremely reconfigurable devices for the home
 - Any-to-any communications and control
 - Standard pluggable interfaces
- **Bioengineering Data Manipulations**
 - Folding@Home
 - Mummer (Genome sequence alignment)
 - mummer.sourceforge.net

Other Possibilities

- Distributed OS-kernel coordination
- Projects from research groups
- Projects from potential employers
- P2P systems
- Social networking
 - FaceBook & Mobility
- Vehicular networks
 - V2V and V2I

Other Possibilities

- Project based on OpenFlow switch
 - see openflowswitch.org
- Project using DCCP (UDP + congestion control)
- DSM where memory pages are shared using a P2P network (challenging, potentially publishable)
- Secure, decentralized, name-system
 - See www.aaronsw.com/weblog/squarezooko